

1. What is the difference between binary and counting semaphores?

A binary semaphore is a counting semaphore that has a count of either 1 or 0. This ensures that only one thread/process will be running at a time. A counting semaphore can have a count of any positive integer and can set a max number of threads running at once

2. What are the operations that can be performed on a semaphore. Explain in brief about each operation and what it does in your own words. A semaphore can either be incremented (semSignal/semPost) or decremented (semWait). Incrementing will increase a semaphore's counter and unblock a process if the semaphore's counter is 0, while decrementing will keep block a process/thread until a different process/thread increments it.

3. **5.6.** Consider the following processes P1 and P2 that update the value of the shared variables, x and y, as follows:

Process P1 : (performs the operations: x := x * y y ++) LOAD R1, X LOAD R2, Y MUL R1, R2 STORE X, R1 INC R2 STORE Y, R2	Process P2 : (performs the operations: x ++ y := x * y) LOAD R3, X INC R3 LOAD R4, Y MUL R4, R3 STORE X, R3 STORE Y, R4
--	--

Assume that the initial values of x and y are 2 and 3 respectively. P1 enters the system first and so it is required that the output is equivalent to a serial execution of P1 followed by P2. The scheduler in the uniprocessor system implements a pseudo-parallel execution of these two concurrent processes by interleaving their instructions without restricting the order of the interleaving.

- If the processes P1 and P2 had executed serially, what would the values of x and y have been after the execution of both processes?
- Write an interleaved concurrent schedule that gives the same output as a serial schedule.
- Write an interleaved concurrent schedule that gives an output that is different from that of a serial schedule.

A. After Process 1 the value of X is 6 and Y is 4. After Process 2, X is 7 and Y is 28.

B. sem_t *semaphore; something to initialize a semaphore
sem_init(sem, 1, 0); start semaphore at 0
Do process 1;
sem_post(sem); signal it's done

sem_wait(sem); make semaphore wait for other process to finish
Do process 2;
sem_destroy;
return 0; finished program

C. sem_t *semaphore; // something to initialize a semaphore
sem_init(sem, 1, 0); start semaphore at 0
Do x:= x*y from process 1;
sem_wait(sem); wait for process 2 to finish
increment X to finish process 1;
Do process 2;
sem_post(sem); to allow process 1 to finish
sem_destroy;
return 0; finished program

4. The following three functions are run on a shared processor by three processes. They can coordinate their execution via shared semaphores that respond to the standard signal(sem_signal()) and wait(sem_wait()) procedures. In order to produce the output HELLO, add respective sem_signal()/sem_post() and sem_wait() comands in the code. Create your own semaphores as needed.

- Is printing HELLO possible
- Number of semaphores - 2
- Names of semaphores - sem1, sem2
- Initial values of sempahores - sem_init(&Sem, 0, 0)

Function#1	Function#2	Function #3
print("H")	sem_wait(&sem1) print("L")	sem_wait(&sem2) print("O")
print("E") sem_post(&sem1)	print("L") sem_post(&sem2)	

終わる